

```
import kagglehub
```

```
# Download latest version
```

```
path = kagglehub.dataset_download("atharvaingle/crop-recommendation-dataset")
```

```
print("Path to dataset files:", path)
```

Downloading from https://www.kaggle.com/api/v1/datasets/download/atharvaingle/crop-recommendation-dataset?dataset_version_number=1...
100%|██████████| 63.7k/63.7k [00:00<00:00, 338kB/s]Extracting files...

```
Path to dataset files: /root/.cache/kagglehub/datasets/atharvaingle/crop-recommendation-dataset/versions/1
```

```
import pandas as pd
```

```
df = pd.read_csv(path + "/Crop_recommendation.csv")
```

```
df.head()
```

```

N    P    K  temperature  humidity      ph  rainfall  label
0  90   42   43    20.879744  82.002744  6.502985  202.935536  rice
1  85   58   41    21.770462  80.319644  7.038096  226.655537  rice
2  60   55   44    23.004459  82.320763  7.840207  263.964248  rice
3  74   35   40    26.491096  80.158363  6.980401  242.864034  rice
4  78   42   42    20.130175  81.604873  7.628473  262.717340  rice

```

```
X = df.drop('label', axis=1) # Features
```

```
y = df['label'] # Labels
```

```
y = pd.get_dummies(y)#instead of being a word output, our model will output a true or false for each type of fruit
```

```
y.head()
```

```

apple  banana  blackgram  chickpea  coconut  coffee  cotton  grapes  jute  kidneybeans  ...  mango  mothbeans  mungbean  muskmelon  or
0  False     False      False     False     False  False  False  False  False      False  ...  False     False     False     False     f
1  False     False      False     False     False  False  False  False  False      False  ...  False     False     False     False     f
2  False     False      False     False     False  False  False  False  False      False  ...  False     False     False     False     f
3  False     False      False     False     False  False  False  False  False      False  ...  False     False     False     False     f
4  False     False      False     False     False  False  False  False  False      False  ...  False     False     False     False     f

```

```
5 rows × 22 columns
```

```
#scales all numerical values from 0 - 1 to make it easier for model to understand data
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
import joblib
```

```
scaler = MinMaxScaler()
```

```
scaler.fit(X)
```

```
X = pd.DataFrame(scaler.transform(X), columns=X.columns)
```

```
joblib.dump(scaler, "/content/drive/MyDrive/Research/TSAFOLDER/saved_scaler.pkl")
```

```
X.head()
```

```

N      P      K  temperature  humidity      ph  rainfall
0  0.642857  0.264286  0.190    0.345886  0.790267  0.466264  0.656458
1  0.607143  0.378571  0.180    0.371445  0.770633  0.549480  0.741675
2  0.428571  0.357143  0.195    0.406854  0.793977  0.674219  0.875710
3  0.528571  0.214286  0.175    0.506901  0.768751  0.540508  0.799905
4  0.557143  0.264286  0.185    0.324378  0.785626  0.641291  0.871231

```

```
#required libraries
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
import tensorflow as tf
```

```
#split out model into training, testing, and validation
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Split the training set into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.25, random_state=42)
```

```
X_train.shape
```

```
(1320, 7)
```

```
import tensorflow as tf
from tensorflow.keras import layers, models, optimizers
from tensorflow.keras.metrics import Precision, Recall

input_dim = 7
# Define the model
model = models.Sequential([
    layers.Input(shape=(input_dim,)), # Replace 'input_dim' with the number of input features
    layers.Dense(128, activation='relu'), # First hidden layer
    layers.Dropout(0.2), # Dropout to reduce overfitting
    layers.Dense(64, activation='relu'), # Second hidden layer
    layers.Dropout(0.2),
    layers.Dense(32, activation='relu'), # Third hidden layer
    layers.Dense(22, activation='softmax') # Output layer (regression task)
])

# Compile the model
model.compile(optimizer=optimizers.Adam(learning_rate=0.001),
              loss='categorical_crossentropy',
              metrics=['accuracy', Precision(name='precision'), Recall(name='recall')])

model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 128)	1,024
dropout_2 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 64)	8,256
dropout_3 (Dropout)	(None, 64)	0
dense_6 (Dense)	(None, 32)	2,080
dense_7 (Dense)	(None, 22)	726

```
Total params: 12,086 (47.21 KB)
Trainable params: 12,086 (47.21 KB)
```

```
# Train the model
history = model.fit(
    X_train, y_train, # Replace with your training data
    validation_data=(X_val, y_val), # Replace with your validation data
    epochs=50, # Adjust based on performance
    batch_size=16, # Suitable for large datasets
    callbacks=[tf.keras.callbacks.EarlyStopping(patience=5, restore_best_weights=True)]
)
```

```
# Evaluate the model on test data
test_loss, test_mae = model.evaluate(X_test, y_test) # Replace with your test data
print(f"Test Loss: {test_loss}, Test accuracy: {test_mae}")
```

```
#gives accuracy of model
test_accracy = model.evaluate(X_test, y_test)
```

```
14/14 ————— 0s 4ms/step - accuracy: 0.9533 - loss: 0.1435 - precision: 0.9542 - recall: 0.9501
```

```
model.save("/content/drive/MyDrive/Research/TSAFOLDER/agriculture-model")
```

```
df.head()
```

	N	P	K	temperature	humidity	ph	rainfall	label
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice

```
#required libraries
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
import numpy as np
```

```
import tensorflow as tf
```

```
#array of possible crops
```

```
crops = [
```

```
"apple", "banana", "blackgram", "chickpea", "coconut", "coffee", "cotton",
"grapes", "jute", "kidneybeans", "lentil", "maize", "mango", "mothbeans",
"mungbean", "muskmelon", "orange", "papaya", "pigeonpeas", "pomegranate",
"rice", "watermelon"
```

```
]
```

```
#scaler to change values from 0 -1
```

```
loaded_scaler = joblib.load("/content/drive/MyDrive/Research/TSAFOLDER/saved_scaler.pkl")
```

```
#load model from paht, you will need to change this
```

```
model = tf.keras.models.load_model("/content/drive/MyDrive/Research/TSAFOLDER/agriculture-model.h5")
```

```
#array of values in the order of nitrogen, phosphorus, potassium, temp, humidity, ph, and rainfall
```

```
values = [74,54,35.4,20.87974371,82.00274423,7,202.9355362]
```

```
values = pd.DataFrame([values],
```

```
columns=['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall'])
```

```
values = pd.DataFrame(loaded_scaler.transform(X), columns=X.columns)
```

```
values = np.array(values)
```

```
values = values.reshape(1,7)
```

```
print(values)
```

```
output = model.predict(values)
```

```
output = np.argmax(output)
```

```
print(output)
```

```
crops[output]
```

```
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you t
[[0.64285714 0.26428571 0.19          0.34588614 0.79026683 0.46626364
 0.65645778]]
```

```
1/1 ----- 0s 230ms/step
```

```
20
```

```
'rice'
```

```
X_test
```

	N	P	K	temperature	humidity	ph	rainfall
1451	0.721429	0.085714	0.210	0.593069	0.938733	0.416821	0.021904
1334	0.700000	0.021429	0.230	0.497956	0.842992	0.428373	0.104973
1761	0.421429	0.407143	0.220	0.990962	0.922659	0.534458	0.339742
1735	0.314286	0.392857	0.250	0.730414	0.890039	0.516399	0.281407
1576	0.214286	0.942857	0.975	0.404267	0.891779	0.326369	0.353489
...
59	0.707143	0.357143	0.150	0.370107	0.769692	0.466063	0.926001
71	0.478571	0.285714	0.165	0.398918	0.792226	0.590274	0.864657
1908	0.864286	0.300000	0.055	0.424105	0.758689	0.656029	0.187846
1958	0.828571	0.335714	0.070	0.405084	0.712913	0.405853	0.168382
482	0.035714	0.450000	0.075	0.293205	0.219879	0.406964	0.485575

```
440 rows × 7 columns
```

Double-click (or enter) to edit

```
#scales all numerical values from 0 - 1 to make it easier for model to understand data
from sklearn.preprocessing import MinMaxScaler
import joblib
scaler = MinMaxScaler()
scaler.fit(X)
```

```
X = pd.DataFrame(scaler.transform(X), columns=X.columns)
joblib.dump(scaler, "/content/drive/MyDrive/Research/TSAFOLDER/saved_scaler.pkl")
```

```
X.head()
```

```
loaded_scaler = joblib.load("/content/drive/MyDrive/Research/TSAFOLDER/saved_scaler.pkl")
```

```
values = [90,42,43,20.87974371,82.00274423,6.502985292000001,202.9355362]
values = pd.DataFrame([values],
                      columns=['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall'])
values = pd.DataFrame(loaded_scaler.transform(X), columns=X.columns)
```

```
values = np.array(values)
values = values.reshape(1,7)
values = loaded_scaler.transform(values)
print(values)
```

```
[[0.64285714 0.26428571 0.19          0.34588614 0.79026683 0.46626364
  0.65645778]]
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but MinMaxSc
warnings.warn(
```

```
X = pd.DataFrame(loaded_scaler.transform(X), columns=X.columns)
X
```

```

   N         P         K  temperature  humidity         ph  rainfall
0  0.642857  0.264286  0.19      0.345886  0.790267  0.466264  0.656458
```